Advanced Descriptor-Based and Adaptive Techniques for 3D Mesh Segmentation

I. Introduction

A. The Imperative of 3D Mesh Segmentation in Modern Computer Graphics

The segmentation of three-dimensional (3D) meshes stands as a cornerstone in contemporary computer graphics and geometric processing. Its fundamental role is to decompose complex 3D models into constituent parts that are meaningful and manageable for subsequent operations.[1] Such partitioning is indispensable across a diverse array of applications, including but not limited to: advanced shape understanding, enabling machines to interpret and analyze geometric forms; reverse engineering, where physical objects are digitized and their components identified; character animation, facilitating the rigging and articulation of digital figures; medical imaging, for isolating anatomical structures or pathological regions [3]; and robotics, where environment perception and object interaction rely on segmented scene representations.[1] The ability to effectively segment meshes underpins progress in these fields by transforming monolithic data into structured, semantically relevant information.[5]

B. User's Context and Objectives

This report acknowledges the user's prior engagement with the field, specifically through the implementation of the "Pairwise Harmonics for Shape Analysis" paper. Building upon this foundation, the user expresses a keen interest in leveraging this implementation, or alternatively, exploring other descriptor-based methodologies for sophisticated 3D mesh segmentation. Furthermore, there is a clear objective to investigate the potential of adaptive mesh segmentation techniques to enhance these processes.

C. Report Scope and Objectives

This document aims to provide an expert-level, technically profound analysis of the "Pairwise Harmonics for Shape Analysis" paper, with a particular focus on its applicability to mesh segmentation. Beyond this specific method, the report will comprehensively survey and critically evaluate a range of prominent alternative descriptor-based techniques. It will also delve into the principles and methodologies of adaptive mesh segmentation strategies. The overarching goal is to furnish the user with a thorough understanding of the current landscape and actionable insights, thereby supporting their research or development activities in the domain of 3D mesh segmentation. The exploration of how different methods compare, especially in light of the user's existing implementation, forms a crucial part of this objective. This involves not just a review of individual techniques, but a comparative analysis that positions "Pairwise Harmonics" within the broader spectrum of segmentation approaches and examines how adaptive strategies can be integrated or prove relevant.

D. Methodological Approach of the Report

The analysis presented herein is synthesized from a careful review of seminal research papers, technical articles, and relevant academic publications. The focus is maintained on elucidating the core algorithmic principles, conducting a comparative assessment of their respective strengths and weaknesses, and examining pertinent computational considerations. This approach ensures a robust and well-grounded exploration of the subject matter.

II. In-Depth Analysis: "Pairwise Harmonics for Shape Analysis"

The paper "Pairwise Harmonics for Shape Analysis" introduces a distinctive approach to understanding and processing 3D shapes, diverging from traditional methods by focusing on relationships between pairs of points rather than analyzing points in isolation.[7]

A. Core Concepts and Theoretical Foundations

The methodology presented in "Pairwise Harmonics for Shape Analysis" [7] introduces a significant shift from conventional point-wise shape analysis techniques. Instead of computing descriptors for individual surface points based on their local neighborhoods, this framework establishes descriptors based on pairs of surface points. This **pairwise analysis paradigm** is central to the method's novelty and efficacy. The rationale is that analyzing pairs of points can offer more global insights, reveal discriminative shape structures, and effectively capture variations in the underlying geometry.[7] Such an approach inherently encodes relational information between two points and the geometric characteristics of the surface connecting them.

At the heart of this pairwise analysis are **harmonic functions** defined on mesh surfaces. A function f is harmonic if it satisfies the Laplace equation ∇2f=0 (or Δf=0), subject to specified Dirichlet boundary conditions.[7] For a given pair of points (p,q) on the mesh, a **pairwise harmonic function**, denoted fp,q, is defined. This function is characterized by the Dirichlet boundary conditions f(p)=0 and f(q)=1. Consequently, the function fp,q establishes a scalar field over the mesh surface with values ranging smoothly from 0 at point p to 1 at point q.[7]

The geometric intuition is that the gradient of such a function is non-zero almost everywhere, and its integral lines represent the "smoothest" paths from p to q. The **isocurves** (or level sets) of this pairwise harmonic function, fp,q=c for constant $c \in $, are then utilized as geometric probes.[7] These isocurves are uniformly sampled from the harmonic field. Close to the source point p and the sink point q, these level sets can be approximated by disks on the surface.[7] Figure 2 in the source paper [7] illustrates a pairwise harmonic field and its uniformly sampled isocurves.

From these isocurves, specific **shape descriptors** are derived to quantify the geometric properties between the point pair (p,q). The primary descriptors introduced are [7]:

1. **Perimeter Descriptor (R):** This descriptor measures the length of the isocurves.
2. **Distance Descriptor (D):** This descriptor quantifies the average distance of the points on an isocurve to the two endpoints p and q.

Histograms of these descriptors can reveal structural information; for instance, Figure 1 in the source paper [7] demonstrates that if points p and q are symmetric, the distributions of both R and D descriptors exhibit near symmetry.

B. Methodology for Shape Analysis and Applications

The pairwise harmonics framework is shown to be applicable to several geometry processing tasks, including intrinsic reflectional symmetry axis computation, matching shape extremities, and, of particular interest here, simultaneous surface segmentation and skeletonization.[7] The method for simultaneous surface segmentation and skeletonization, detailed in Section 3.3 of the original paper and further elaborated in [7], leverages the dual nature of the isocurves derived from pairwise harmonic fields. These isocurves can be interpreted both as cross-sectional profiles (useful for skeleton extraction) and as potential boundaries for segmenting the mesh into meaningful parts.[7]

The algorithm for simultaneous segmentation and skeletonization proceeds through the following key steps [7]:

1. **Surface Sampling:** A max-min sampling strategy is employed to select a sparse set of point pairs. This strategy aims to distribute the pairs such that the generated pairwise harmonic fields cover the entire model.
2. **Pairwise Harmonic Computation:** For each sampled pair of points (p,q), the corresponding pairwise harmonic function fp,q is computed.
3. **Skeletal Segment Generation:** The centroids of the isocurves extracted from each pairwise harmonic field are computed and connected sequentially to form initial skeletal segments.
4. **Node Rigidity Assignment:** Each skeletal node (isocurve centroid) is assigned a rigidity score. This score is derived from the eigenvalues of the principal components computed on a local neighborhood of the node. Nodes with rigidity below a certain threshold (e.g., < 0.9) are classified as nonrigid.
5. **Skeletal Segment Division:** The initial skeletal segments are divided into subsegments at the identified nonrigid nodes. This step is crucial for identifying potential articulation points.
6. **Segment Scoring and Selection:** Each resulting subsegment is assigned a score based on its cumulative rigidity and length. A greedy algorithm is then used to select a set of disconnected, prominent skeletal segments that have the highest scores. These selected segments form a partial skeleton of the shape.
7. **Initial Segmentation:** The mesh is partitioned into initial segments using the isocurves

associated with the two end nodes (centroids corresponding to isocurves near p and q) of each selected skeletal segment from the partial skeleton.

8. **Skeleton Completion:** For regions of the mesh that form junctions between multiple initial segments, center nodes are computed. These junction nodes are then connected to adjacent skeletal segments based on the adjacency of the segmentation regions, thereby completing the skeleton.
9. **Segmentation Refinement:** The initial segmentation is further refined through a series of merging operations:
   - Components corresponding to skeletal endpoints are merged with their sole adjacent component.
   - For skeletal segments that are incident to two junction nodes, the mesh component associated with this segment is iteratively merged with its two adjacent components based on component volume, until no further merging is possible. Junction nodes involved in such merges are tagged to prevent re-processing.

C. Strengths and Advantages for Mesh Segmentation

The pairwise harmonics approach offers several compelling advantages for mesh segmentation tasks:

- **Global and Discriminative Analysis:** By defining descriptors based on point pairs, the method inherently captures not only the geometry between the points but also their relationship. This facilitates a more global understanding of the shape and is more discriminative in revealing underlying structures, such as symmetries, compared to traditional descriptors defined solely on local neighborhoods of single points.[7]
- **Inherent Multiscale Analysis:** The pairwise analysis framework naturally incorporates multiscale shape information without requiring additional computational overhead. Measurements are constrained along local geodesic paths between point pairs, allowing the method to reveal shape structures at different scales for different pairs within a single, unified framework.[7]
- **Robustness:** A significant advantage is the robustness of the derived descriptors. Because harmonic functions are smooth and depend on global boundary conditions, the resulting perimeter (R) and distance (D) descriptors are relatively insensitive to variations in mesh tessellation, surface noise, and changes in pose (isometric deformations).[7] This is clearly illustrated in Figure 2 of the source paper.[7]
- **Alignment with Rigidity and Articulatory Structure:** The simultaneous segmentation and skeletonization method is designed to produce segmentations that align well with the perceived rigidity of shape components. The use of a "node rigidity" measure helps identify rigid parts and branching structures, leading to segmentations that can reflect the articulatory structure of an object. This is a key distinction from many other segmentation methods that may not explicitly recover such information.[7]
- **Meaningful Skeletons:** The skeletons generated by this method tend to resemble a rigid decomposition of the shape, with junction nodes correctly located at branching

surfaces and each skeletal branch corresponding to a distinct rigid component of the shape.[7]

- **Computational Efficiency (Claimed):** The authors assert that their approach leads to simpler and more efficient algorithms compared to some state-of-the-art methods, such as those based on mesh contraction. For the segmentation task, computation times of under 30 seconds for typical models are reported.[7]

D. Limitations and Potential Challenges

Despite its strengths, the pairwise harmonics method for segmentation is not without limitations:

- **Branch Identification Dependency:** The algorithm's ability to identify skeletal branches, and consequently to segment parts correctly, can be compromised if a shape component does not possess a well-defined principal axis. In such cases, a corresponding skeletal branch might not be identified, leading to under-segmentation. An example cited is the failure to separate the body and head of a teddy bear model.[7]
- **Exclusion of Certain Model Categories in Evaluation:** The quantitative evaluation presented for the segmentation and skeletonization method in [7] notably excluded certain categories of models, such as heads, vases, fish, and CAD models. The reason provided was that the skeletons for these shapes are not well-defined, which suggests potential limitations in the method's applicability or robustness when dealing with such object types.[7]
- **Computational Cost of Harmonic Functions:** While overall efficiency is claimed, the computation of numerous pairwise harmonic functions can still be computationally intensive, particularly for very dense meshes or if a large number of point pairs are sampled. The practical efficiency likely relies heavily on the use of fast and optimized solvers for the underlying Laplace equations (e.g., methods exploiting sparse matrix structures or fast updates like the Cholesky factorization mentioned in [7]). The pairwise nature introduces a combinatorial aspect, as a naive consideration of all possible pairs of N vertices would be O(N2). The reliance on *sparse sampling* is therefore critical. The effectiveness of this sparse sampling—how well it captures the essential structural relationships needed for segmentation without incurring prohibitive computational costs—is a key factor. If crucial relationships are missed due to overly sparse sampling, the quality of the segmentation could be adversely affected. This implies that the sampling strategy itself is a critical, and potentially sensitive, component of the method's success.
- **Implicit Definition of "Meaningful Parts":** The method's reliance on "node rigidity" and the concept of a "principal axis" for guiding segmentation [7] suggests an implicit assumption about what constitutes a "meaningful part." Such parts are often expected to be somewhat elongated or to possess a dominant axis, with transitions between parts frequently occurring at non-rigid or flexible areas. This paradigm is highly effective for segmenting articulated objects into their "rigid" components, which aligns with a mechanical or articulatory interpretation of shape. However, this approach might

be less effective for segmenting objects based on other criteria, such as segmenting a cube into its faces or partitioning an object based on material properties (if such data were available). Thus, the user must consider whether this implicit definition of a "part" aligns with the specific goals of their segmentation task.

E. Computational Aspects and Implementation Notes for Segmentation

For effectively implementing the segmentation aspect of the pairwise harmonics method, several computational elements are critical:

- **Sparse Sampling Strategy:** The use of a max-min sampling strategy is fundamental to managing computational load by limiting the number of pairwise harmonic fields that need to be computed.[7] The quality of this sampling directly impacts coverage and the ability to capture relevant shape structures.
- **Harmonic Function Solvers:** The core computational task is solving the Laplace equation $\Delta f=0$ with Dirichlet boundary conditions $f(p)=0, f(q)=1$ on the mesh. This is typically formulated as a sparse linear system. Efficient numerical methods, such as the Finite Element Method (FEM) or Finite Difference Method adapted for meshes, are required. The efficiency gains mentioned in [7] likely stem from using advanced techniques like preconditioned conjugate gradient solvers or direct solvers that can exploit matrix sparsity and structure, such as fast Cholesky factorization updating for related problems.
- **Isocurve Extraction and Centroid Computation:** Once the harmonic field $f_{p,q}$ is computed (i.e., scalar values at each vertex), algorithms are needed to efficiently trace the isocurves (level sets). This can involve linear interpolation along mesh edges. Subsequently, the geometric properties of these isocurves, such as their lengths (for the R descriptor) and the coordinates of their centroids (for skeleton generation), must be computed.
- **Greedy Algorithms:** The selection of prominent skeletal segments and the subsequent segmentation refinement steps utilize greedy algorithms.[7] While computationally efficient, greedy approaches do not guarantee globally optimal solutions. The order of operations or the specific scoring functions can influence the final segmentation.

F. Suitability and Application of the User's Implementation for Mesh Segmentation Tasks

The user's existing implementation of "Pairwise Harmonics for Shape Analysis" can be directly focused on mesh segmentation by implementing or adapting the simultaneous segmentation and skeletonization algorithm detailed in.[7] Key aspects to consider for such an application include:

- **Parameter Tuning:** Successful segmentation will likely require careful tuning of several parameters:
  - *Sampling Density:* The number of point pairs selected by the max-min strategy. Too few may miss important features; too many will increase computation time.

- *Rigidity Thresholds:* The threshold used to classify skeletal nodes as rigid or nonrigid. This directly affects where skeletal segments are divided and, consequently, where initial segment boundaries are placed.
- *Merging Criteria:* Parameters controlling the segmentation refinement process, such as volume thresholds or adjacency rules for merging components.

- **Strengths and Weaknesses in Specific Scenarios:**
  - The method is expected to perform particularly well for segmenting **articulated objects** (e.g., humans, animals) where parts are connected by flexible joints, as the rigidity analysis is designed to capture this.
  - Shapes with clear **symmetries** or **elongated parts** are also good candidates, as pairwise analysis is adept at revealing such global structures.
  - It might struggle with **blobby objects** that lack clear axes or distinct rigid components, or objects where meaningful parts are not defined by articulations (e.g., segmenting a car into hood, doors, trunk if these are not primarily defined by rigidity changes but by feature lines).
  - As noted, objects without well-defined **principal axes** for their components can pose a challenge.[7]

The user should evaluate their implementation on datasets that play to these strengths and be mindful of the potential limitations when applying it to other types of shapes.

III. A Survey of Prominent Descriptor Methods for 3D Mesh Segmentation

Descriptor-based mesh segmentation is a well-established paradigm in 3D shape analysis. The general principle involves deriving a feature vector, or descriptor, for each element of the mesh (such as vertices or faces) or for local regions. These descriptors aim to capture salient geometric or topological properties. Once computed, these descriptors are then used as input to clustering algorithms or classification frameworks to group mesh elements into distinct, meaningful segments.[8] Early approaches often relied on hand-crafted geometric features [8], while more recent work has explored a wide variety of sophisticated descriptors, including those derived from spectral analysis, volumetric properties, and topological abstractions.

A. Spectral Descriptors

Spectral mesh processing offers a powerful framework for analyzing 3D shapes by "porting the signal processing toolbox" to the domain of meshes.[10] This approach leverages the eigenstructures of operators defined on the mesh, most notably the mesh Laplacian.
1. Theoretical Foundations:
The core idea is to analyze the mesh in a "frequency domain" defined by the eigenfunctions of a mesh Laplacian operator.

- **Mesh Laplacian Operators:** These operators are central to spectral mesh analysis and act on functions defined over the mesh vertices or faces. They capture information about both the connectivity (topology) and, in some forms, the geometry of the mesh.[10] Two main categories exist:

  - **Combinatorial Laplacians:** These are defined solely based on the mesh connectivity. A common example is the graph Laplacian, often defined as $K=D-W$, where D is the diagonal matrix of vertex degrees and W is the adjacency matrix of the mesh graph.[12] These Laplacians treat the mesh as an abstract graph.
  - **Geometric Laplacians:** These operators are discrete approximations of the continuous Laplace-Beltrami operator from Riemannian geometry and explicitly incorporate geometric information, such as edge lengths and angles. The most widely used is the **cotangent Laplacian** (or cotan Laplacian). For a function f defined at vertices, its value at vertex i is typically given by $[L_c f]_i=\sum_{j\in N(i)} w_{ij}(f_j-f_i)$, where the weights $w_{ij}=\frac{1}{2}(\cot\alpha_{ij}+\cot\beta_{ij})$, and α_{ij},β_{ij} are the angles opposite to the edge (i,j) in the two triangles sharing that edge.[11] Geometric Laplacians are generally preferred for shape analysis tasks as they more faithfully reflect the intrinsic geometry of the surface.
- **Eigen-decomposition:** The mesh Laplacian, represented as a matrix M, is subjected to an eigen-decomposition to find its eigenvalues λ and corresponding eigenvectors v, satisfying the equation $Mv=\lambda v$.[10]
  - **Eigenvalues (λ):** These are real, non-negative values often interpreted as the "frequencies" of the mesh. They provide global information about the shape's structure. Small eigenvalues correspond to low frequencies (smooth, global variations), while large eigenvalues correspond to high frequencies (sharp, local details).[11]
  - **Eigenvectors (v):** Also known as eigenfunctions, these form an orthogonal basis for functions defined on the mesh. Each eigenvector represents a fundamental "mode of vibration" or variation pattern of the shape. The eigenvectors corresponding to the smallest non-zero eigenvalues (low-frequency modes) capture the global, smooth structure of the shape and can be used to create a **spectral embedding** of the mesh into a lower-dimensional Euclidean space.[11]

2. Segmentation Techniques using Spectral Descriptors:

The eigenstructures of the mesh Laplacian provide rich information for segmentation:

- **Spectral Clustering:** This is a common technique where a few of the lowest-frequency eigenvectors are used to embed the mesh vertices (or faces) into a low-dimensional spectral space. In this space, geometrically or topologically related elements tend to cluster together. Standard clustering algorithms, such as k-means, are then applied to these embedded points to partition the mesh.[12]
- **Recursive Spectral Cuts and Nodal Domains:** Eigenvectors can directly guide the partitioning of a mesh. The **Fiedler vector**, which is the eigenvector corresponding to the second smallest eigenvalue of the graph Laplacian, is particularly known for its use in graph bipartitioning.[13] Its sign or thresholding its values can effectively divide a mesh into two parts. This process can be applied recursively to obtain more segments. More generally, the **nodal domains** of an eigenfunction (regions where the eigenfunction

maintains a consistent sign, separated by zero-crossings or nodal lines) can define natural segments on the mesh.[12]

- **"Novel Single Segmentation Field" [13]:** This method, presented by Wang et al., introduces a hierarchical spectral analysis approach using a "concavity-aware Laplacian." The key idea is to modify the standard Laplacian to be more sensitive to concave regions, which often delineate meaningful part boundaries. The process involves:
    1. Constructing this concavity-aware Laplacian.
    2. Adaptively selecting a sufficient number of its eigenvectors.
    3. Partitioning these eigenvectors into sub-eigenvectors using spectral clustering.
    4. Evaluating the confidence of each sub-eigenvector for identifying spectral-sensitive mesh boundaries based on "inner variations" and "part oscillations."
    5. Formulating an optimization problem to select and combine these sub-eigenvectors into a single scalar segmentation field defined over the mesh.
    6. Detecting segmentation boundaries from this single field using a divide-merge algorithm and a "cut score." The concavity-aware Laplacian attempts to make the abstract spectral information more directly relevant to perceptual segmentation by incorporating the heuristic that concave boundaries are important. This aims to bridge the gap between the mathematical elegance of spectral analysis and the practical requirements of "meaningful" segmentation.
- **Farthest Sampling Segmentation (FSS) [13]:** While not strictly a spectral method in the sense of eigendecomposition, FSS is often discussed in relation to spectral clustering. It operates by:
    1. Selecting a sparse sample of k faces from the mesh using farthest point sampling with respect to a chosen metric.
    2. Computing an n×k affinity matrix Wk, where n is the total number of faces. Each row of Wk encodes the similarity (based on pairwise distances) of a mesh face to the k sampled faces.
    3. Applying the k-means++ clustering algorithm to the rows of Wk to obtain the segmentation. FSS is noted for being parameter-free (except for k, the number of segments), flexible due to its ability to handle any metric, and computationally efficient due to the sparse sampling.[13]

A common challenge in many spectral and clustering-based segmentation methods is the determination of the number of segments, k. While some methods, like the "novel single segmentation field" approach, aim to derive boundaries from a continuous field, potentially sidestepping a direct pre-specification of k for the final output, the choice of k (or the number of eigenvectors to retain) often remains a critical parameter that may require user input or sophisticated heuristics.[13] Incorrect choices can easily lead to under- or over-segmentation. Some research does explore automatic determination mechanisms for the number of segments.[13]

B. Volumetric and Point-Based Descriptors (Intrinsic, Robust to

Deformations)

This category includes descriptors that aim to capture intrinsic properties of the shape, often with a focus on robustness to various transformations like pose changes or non-rigid deformations.

1. Shape Diameter Function (SDF):

- **Principle:** The SDF is a scalar function defined at each point on the surface of a closed 3D manifold. It measures an approximation of the local "thickness" or diameter of the object in the neighborhood of that point, thereby capturing local volumetric characteristics.[14] While related to the medial axis transform, SDF is generally considered more robust to compute, especially in the presence of small boundary perturbations.[14]
- **Computation:**
    - *Original Method (Gal et al.):* For a point p on the surface, an inward-pointing cone is conceptually placed at p. Multiple rays are cast within this cone, and their penetration distances through the object are measured. The SDF value at p is then approximated as a weighted average of these penetration distances.[14]
    - *Robust/Offset Method:* To address the sensitivities of the original method, an improved approach based on an offset surface was proposed.[14] The input surface S is first offset outwards by a small distance r to create an offset surface S′. For a point p on the original surface S, its corresponding point q on S′ (along the outward normal) is found. A ray is cast from q through p, and its second intersection point q′ with the offset surface S′ is determined. The SDF value at p is then computed as the distance $// qq' // -2r$. This computation can be made efficient using an Oriented Bounding Box (OBB) tree constructed from basic geometric elements that approximate the offset volume.[14]
- **Properties:**
    - *Pose Oblivious:* SDF values are inherently invariant to rigid transformations (translation and rotation) of the object.[14]
    - *Original SDF Sensitivities:* The original SDF computation is sensitive to fine geometric details, surface noise, and inaccuracies in normal vector estimation. It is also computationally expensive due to the multiple ray casts per point.[14]
    - *Offset SDF Robustness:* The offset-based method significantly improves robustness. It is less sensitive to noise and small geometric details, can handle meshes with gaps or even point clouds (if r is chosen appropriately), and is computationally much faster (reportedly about 10 times) than the original method because it effectively requires a single penetration distance calculation per point.[14]
- **Segmentation Utility:** SDF provides a scalar value at each point reflecting local thickness. This field of SDF values can be used directly for segmentation, for instance, by clustering vertices or faces that have similar SDF values. Histograms of SDF values can also be used as features to characterize parts.[14] The offset-based SDF has shown improved segmentation performance over the original method.[14]

2. **Heat Kernel Signature (HKS):**
   - **Principle:** HKS is a point-wise shape descriptor based on the principles of heat diffusion on a Riemannian manifold (the mesh surface). It characterizes the amount of heat remaining at a point after a certain diffusion time, starting from an initial impulse at that same point. HKS is multi-scale, intrinsic to the shape, and invariant to isometric deformations.[3]
   - **Computation:** The HKS is derived from the heat kernel $k_t(x,y)$, which describes the heat flow from point y to point x in time t. The heat kernel can be expressed using the eigenvalues λ_i and eigenfunctions ϕ_i of the Laplace-Beltrami operator on the mesh: $k_t(x,y)=\sum_{i=0}^{\infty} e^{-\lambda_i t}\phi_i(x)\phi_i(y)$. The HKS at a point x is then defined as a vector of diagonal values of the heat kernel for a set of discrete time scales $t_1,t_2,\dots,t_n$: $HKS(x)=(k_{t_1}(x,x),k_{t_2}(x,x),\dots,k_{t_n}(x,x))$.[16] The time parameter t acts as a scale parameter: small t values capture local, fine details, while larger t values reflect more global geometric properties.
   - **Segmentation Utility:**
     - HKS serves as a multi-scale signature of local curvature and geometry. The local maxima of the HKS function (for a fixed t) can be used to identify salient feature points on the mesh.
     - A notable application combines HKS with **Persistence-Based Clustering (PBC)**. Here, the HKS function $f(x)=k_t(x,x)$ is defined over the mesh. PBC analyzes the topological evolution of the superlevel sets of this function. By examining the "persistence" (lifetime) of connected components as the function threshold varies, stable segments are identified. A merging parameter τ, derived from the persistence diagram, controls the granularity of the resulting segmentation. This combination yields a stable, isometry-invariant segmentation method.
     - HKS can also be used to derive global shape descriptors (e.g., HeatSD, a histogram-based descriptor from HKS values [16]), which, while not directly segmenting, can inform segmentation if shapes are first classified or matched based on these global descriptors.
3. **Wave Kernel Signature (WKS):**
   - **Principle:** WKS is another point-wise spectral descriptor, analogous to HKS, but derived from the time-dependent Schrödinger wave equation on the manifold. It is also designed to be robust to isometric deformations and capture multi-scale geometric information.[17]
   - **Computation (Improved WKS - IWKS [17]):**
     - The standard WKS for a point x and energy scale E is $WKS(E,x)=\sum_{k=1}^{\infty}\phi_k(x)^2 f_E(\lambda_k)^2$, where $f_E(\lambda_k)$ is a band-pass filter (often log-normal) centered at energy E and applied to eigenvalue λ_k.
     - The **Improved WKS (IWKS)** introduces two key modifications to enhance discriminative power and robustness for non-rigid shapes:
       1. **Adaptive Eigenvalue Scaling:** Instead of logarithmic scaling of eigenvalues for the filter, IWKS employs a **cubic root scaling**. The filter

$f_C(\Lambda_k,e)^2=\exp(-2\sigma^2(e-\Lambda_k)^2)$ (where Λ_k is the eigenvalue and e is the energy scale) is found to distribute the band-pass filters more evenly across the spectrum, capturing both coarse and fine details more effectively.

2. **Curvature Aggregation:** IWKS incorporates local shape curvature by adding a term based on the maximum principal curvature C(x) at each point: $IWKS(x,e)=\sum_{k=1}^{\infty}\phi_k(x)^2 f_C(\Lambda_k,e)^2+\alpha C(x)$. This helps to downweight regions prone to articulation (often low curvature) and emphasize more stable, geometrically significant parts.

- **Segmentation Utility:** IWKS provides a discriminative local descriptor. Similar to HKS, these point-wise descriptors can be clustered, or used as features in a machine learning framework, to achieve mesh segmentation. Its design for robust non-rigid shape retrieval suggests its utility for segmenting deformable objects or objects across different poses.[17] The general approach of using feature vectors (which IWKS provides per point) for clustering to assign labels is a common segmentation strategy.[8]

The development of descriptors like SDF, HKS, and WKS underscores a drive to encode rich, intrinsic local geometry that remains stable under various transformations. The multi-scale nature of HKS and WKS, achieved through the time or energy parameters, is particularly valuable. This allows these descriptors to capture features at different levels of detail, which is essential for hierarchical segmentation or for understanding context in complex shapes. Furthermore, the evolution from original SDF to the offset-SDF, and from WKS to IWKS, illustrates a common pattern in descriptor research. Initial theoretically grounded descriptors are often refined through practical application to enhance their robustness against real-world data imperfections (like noise and gaps) and to improve their discriminative power for specific challenges (like non-rigid deformations). This often involves incorporating additional geometric heuristics (e.g., curvature weighting in IWKS) or devising new computational strategies (e.g., the offset surface for SDF). This iterative cycle of theoretical proposal, practical testing, and pragmatic refinement is driven by the complexities of real-world 3D data.

C. Geometric Saliency and Curvature-Driven Methods

These methods leverage local geometric properties, primarily curvature and related saliency measures, to identify distinct regions or boundaries on a mesh surface.[2] The underlying assumption is that perceptually salient features—such as regions of high curvature, sharp edges, or prominent protrusions—often correspond to the boundaries of meaningful parts.

1. Principles:

The core idea is that changes in local surface geometry can indicate segment boundaries. For instance, the transition from a relatively flat area to a sharply curved one might signify the start of a new part. Geometric saliency aims to quantify the "importance" or "noticeableness" of different points or regions on the mesh based on their local shape characteristics. One example describes using a combination of Gaussian-weighted curvature and spin-image correlation to measure geometric saliency for 3D facial mesh detection, where salient vertices with similar properties are subsequently clustered.4

2. Techniques:

- **Curvature-based Segmentation:** This is a direct approach, exemplified by the pebble segmentation method described in.[19] The typical workflow involves:
    1. Computing a curvature measure (e.g., mean curvature, Gaussian curvature, principal curvatures) at each vertex or face of the mesh.
    2. Using curvature values to distinguish between different types of surface regions. For example, convex regions (positive curvature) might be considered parts of objects, while concave regions (negative curvature) might indicate junctions or depressions between parts.
    3. Grouping connected components of mesh elements that share similar curvature properties (e.g., all connected convex triangles).
    4. Applying post-processing steps such as morphological operations (erosion, dilation), hole filling, and filtering based on the shape properties of the resulting components (e.g., sphericity, size) to refine the segmentation.
- **Region Growing:** These algorithms start from one or more initial "seed" points or regions on the mesh. The regions then iteratively expand by annexing neighboring mesh elements (vertices or faces) as long as these neighbors satisfy a homogeneity criterion based on some geometric property (e.g., similar curvature, normal orientation, or planarity).[5]
- **Boundary Detection:** Instead of growing regions, these methods focus on directly identifying the boundaries between segments. This often involves searching for paths of vertices or edges that exhibit high curvature, sharp dihedral angles, or other discontinuities that suggest a transition between parts.[5]

3. Challenges:

- **Sensitivity to Noise and Mesh Irregularities:** Curvature is typically computed using second-order derivatives of the surface. This makes curvature-based methods highly sensitive to noise and variations in mesh tessellation, which can lead to unreliable curvature estimates and spurious segmentations.[2] Pre-smoothing the mesh is often a necessary but delicate preprocessing step.
- **Difficulty with Global Context:** Purely local saliency or curvature measures can struggle to capture global shape context or to segment objects that have smooth transitions between parts. They may over-segment regions with rich local detail or fail to identify boundaries that are subtle locally but significant globally.[2] The success of the pebble segmentation example [20] is partly due to the constrained domain where pebbles are generally convex and distinct. Generalizing this to complex, articulated shapes with varied part definitions based solely on local curvature is more challenging. While intuitive and often computationally less expensive for initial feature detection than global methods, their locality suggests they might be best used as part of a hierarchical approach or combined with methods that provide more global context.

D. Topological Descriptors: Reeb Graphs

Reeb graphs offer a way to analyze and represent the topological structure of a shape by examining the evolution of the connected components of the level sets of a scalar function defined on its surface.[2]

1. Principles:

Given a manifold (the mesh surface) and a continuous scalar function f defined on it (e.g., a height function, geodesic distance from a source point, a curvature measure, or even HKS values), the Reeb graph tracks how the topology of the level sets f−1(c) changes as the value c varies. Nodes in the Reeb graph correspond to the critical points of the function f (minima, maxima, and saddle points), where the topology of the level sets changes. Arcs in the Reeb graph connect these critical points and represent continuous "sheets" or "tunnels" of the manifold between them.

2. Construction:

The construction process typically involves 2:

1. **Defining a Scalar Function:** A suitable scalar function is chosen and evaluated at all vertices of the mesh. The choice of this function is crucial as it determines the resulting Reeb graph and the aspects of the shape's topology that will be captured.
2. **Identifying Critical Points:** The critical points of this scalar function (local minima, local maxima, and saddle points) are identified on the mesh. Morse theory provides the theoretical underpinning for this step.[2]
3. **Graph Construction:** The Reeb graph is constructed by creating nodes for these critical points and adding arcs between them based on the connectivity of the regions between corresponding level sets. An "enhanced topological skeleton approach" is mentioned in [21] for efficient construction that also preserves degenerate critical points.

The choice of the scalar function is paramount in Reeb graph methods. Different functions will highlight different topological aspects of the shape, leading to different Reeb graphs and, consequently, different segmentations. For instance, a height function might capture protrusions along a specific axis, while a geodesic distance function might reveal concentric structures. This provides flexibility but also implies that the "correct" scalar function is application-dependent and there is no universally optimal choice.

3. Topological Simplification:

Reeb graphs derived from complex shapes or noisy scalar functions can be overly detailed and contain many small, insignificant topological features. Therefore, a simplification step is often necessary.21 This typically involves canceling pairs of critical points (e.g., a saddle and an adjacent extremum) that represent minor topological variations, thereby reducing the graph's complexity while aiming to preserve its essential structural features.

4. Segmentation Utility:

The Reeb graph provides a compact, high-level abstraction of the shape's topological structure, often resembling a skeleton. This graph can then be used to guide the segmentation of the mesh:

- Segments on the mesh can be generated by mapping regions of the mesh back from the arcs or subtrees of the simplified Reeb graph. Each arc, for example, might correspond to a limb-like part of an object.

- **Adaptive Region Growing [2]:** A challenge with direct Reeb graph segmentation is that mapping graph components back to the mesh can sometimes result in disconnected regions on the mesh surface, particularly for arcs connecting two saddle nodes or when the scalar function is noisy. To address this, an adaptive region-growing process can be employed. This process integrates both the topological adjacency information from the Reeb graph (guiding which parts should belong together conceptually) and the geometric connectivity on the mesh (ensuring that the resulting segments are contiguous on the surface). Typically, regions are grown from nodes corresponding to extrema, with a priority queue managing candidate nodes for addition to a region based on both graph adjacency and mesh connectivity. This hybrid approach, combining topological guidance with geometric constraints, is crucial for bridging the gap between the abstract Reeb graph and a practical, well-formed segmentation on the concrete mesh. It tends to produce more robust and visually appealing segmentations than relying solely on either topology or geometry.

E. Other Pairwise Descriptor Approaches

Beyond Pairwise Harmonics, other methods also adopt a pairwise analysis philosophy.

1. Pairwise 3D Shape Context [23]:

- **Principle:** This descriptor is defined for pairs of points (A,B) on a 3D shape (often a point cloud in the original application). It aims to simultaneously model the local topology around each point and the global geometric structure of the shape as it exists between these two points, typically within a manifold space.
- **Computation:**
    1. For a given pair of points (A,B), a Region of Interest (ROI) is determined. This might involve finding the shortest path between A and B in manifold space.
    2. The ROI is then partitioned into a set of bins, often using a log-polar or equal interval structure along the radial direction from one of the points.
    3. The distribution of other points from the shape that fall within each bin is analyzed and normalized to form a histogram descriptor for the pair (A,B).
- **Properties:** The resulting pairwise 3D shape context is designed to be invariant to scale and orientation, and partially insensitive to topological changes (e.g., small holes or noise).
- **Application:** The method was originally developed and demonstrated for tasks like partial object matching and retrieval, specifically for extracting 3D light poles and trees from Mobile Laser Scanning (MLS) point clouds.[23]
- **Relevance to Segmentation:** Although showcased for matching, the descriptor inherently captures rich pairwise geometric relationships. If computed for various pairs of points *within* a single complex shape, these descriptors could potentially be used to define similarity measures between different regions or to cluster points/faces for segmentation. This approach shares the spirit of Pairwise Harmonics in focusing on point pairs but uses a distinctly different mechanism (binned point distributions along paths rather than harmonic fields and isocurves) to formulate the descriptor.

The existence of multiple "pairwise" descriptor methods, such as Pairwise Harmonics [7] and Pairwise 3D Shape Context [23], highlights a recognized value in analyzing the relationships *between* points, rather than solely focusing on properties *at* individual points. This relational perspective can capture more global context. However, these methods differ significantly in their underlying mathematical tools (harmonic functions and PDE solutions versus shortest path computations and histogram binning) and therefore in how they define and compute these pairwise features. This implies that they will likely exhibit different sensitivities to various geometric properties, different robustness characteristics (e.g., harmonic functions are inherently smooth, while shortest paths can be more sensitive to mesh tessellation), and different computational profiles.

F. Comparative Table of Descriptor Families

To provide a concise overview, the following table summarizes the key characteristics of the discussed descriptor families relevant to mesh segmentation.

Descriptor Family	Core Principle	Key Features/Properties	Strengths for Segmentation	Weaknesses/Limitations	Typical Computational Cost	Notable References
Pairwise Harmonics	Isocurves of pairwise harmonic fields (f(p)=0,f(q)=1)	Global pairwise view, multiscale, robust to noise/pose	Good for articulated parts, symmetry detection, rigidity-based segmentation	Principal axis dependency for some parts, cost for many pairs, sampling strategy is critical	Moderate to High	[7]
Spectral (Laplacian)	Eigenstructures of mesh Laplacian (combinatorial or geometric)	Intrinsic, global, captures modes of variation, basis for functions on mesh	Captures global structure, good for consistent cuts, basis for spectral clustering, concavity-aware variants exist	Choice of k (segments/eigenvectors) often crucial, features can be abstract, eigendecomposition is costly for large meshes	High	[10]
Shape Diameter Function	Local volumetric thickness/dia	Pose-oblivious, volumetric,	Good for thickness-based parts,	Original SDF sensitive to noise/details;	Low to Moderate (offset)	[14]

(SDF)	meter measurement	robust (offset method)	robust to noise (offset version)	offset SDF may lose fine details; requires closed manifold		
Heat Kernel Signature (HKS)	Point-wise heat diffusion signature over time scales	Isometry-invariant, multiscale, intrinsic	Handles non-rigid shapes, multi-scale feature identification, can find feature points, stable with PBC	Computationally intensive (eigendecomposition needed), parameter tuning (t-scales)	High	[3]
Wave Kernel Signature (WKS)	Point-wise Schrödinger wave equation solution over energy scales	Isometry-invariant, multiscale, intrinsic, IWKS variant more discriminative	Good for non-rigid shapes, discriminative local features (IWKS), multi-scale analysis	Computationally intensive (eigendecomposition needed), parameter tuning (e-scales, σ)	High	[17]
Geometric Saliency/Curvature	Local surface curvature or other geometric distinctiveness measures	Local, intuitive, often computationally inexpensive for basic measures	Fast for identifying sharp features, edges, convex/concave regions	Highly sensitive to noise and tessellation, local view may miss global context, struggles with smooth transitions	Low	[18]
Reeb Graphs	Topology of level sets of a scalar function on the mesh	Topological, skeletal, abstracts shape structure	Defines high-level part structure, good for	Choice of scalar function is critical and application-	Moderate to High	[2]

			shapes with clear topological components (limbs, handles)	dependent, graph construction /simplification can be complex		
Pairwise 3D Shape Context	Binned point distributions along paths between point pairs in manifold space	Pairwise, invariant to scale/orientation, partially insensitive to topological changes	Captures relational context between points, robust to some variations	Primarily demonstrated for matching/retrieval, applicability to dense segmentation needs more exploration, path finding cost	Moderate	[23]

This table serves as a comparative guide, allowing for a quick assessment of which descriptor families might be most suitable for a given segmentation task, considering the nature of the 3D data, the desired properties of the segments, and available computational resources.

IV. Adaptive Mesh Segmentation: Principles and Advanced Techniques

Adaptive techniques in mesh processing aim to dynamically adjust the mesh structure or the computational process itself to improve accuracy and/or efficiency, typically by concentrating effort in regions where it is most needed. While Adaptive Mesh Refinement (AMR) is well-established in numerical simulations for solving PDEs [24], its principles can be extended and adapted for the task of 3D mesh segmentation.

A. Fundamentals of Adaptive Meshing for Segmentation

Concept Definition:
In the context of numerical analysis, Adaptive Mesh Refinement (AMR) is a method that dynamically alters the computational grid's resolution during a simulation to enhance the accuracy of the solution in specific "sensitive" or "turbulent" regions.24 For 3D mesh segmentation, this concept translates to modifying the mesh structure (e.g., vertex density, element shapes) or the segmentation algorithm's behavior based on evolving geometric features, descriptor values, or the confidence in current segment boundaries. The primary goal shifts from PDE solution accuracy to improving the quality (e.g., boundary precision, semantic correctness) and efficiency of the segmentation process.

Motivation:
Uniformly meshed models can be suboptimal; manual creation might lead to excessive cell counts (over-refinement) in simple areas or insufficient resolution (under-refinement) in detailed regions, impacting both solve times and accuracy.25 Adaptive methods strive to create an optimal mesh by selectively focusing computational resources and mesh density on critical areas relevant to the segmentation task.25
Adaptation Criteria for Segmentation:
The "intelligence" in adaptive segmentation lies in defining criteria that effectively guide the adaptation process towards better segmentation outcomes. These criteria can be inspired by AMR principles but must be tailored for segmentation goals:

- **Error Estimators / Segmentation Confidence:** Analogous to error estimators like Richardson extrapolation used in AMR for PDE solutions [1], one could define criteria based on:
  - *Segmentation uncertainty:* Regions where the classification confidence of mesh elements is low.
  - *Boundary stability:* Areas where segment boundaries fluctuate significantly under small perturbations or with different algorithm parameters.
  - *Inter-descriptor agreement:* Discrepancies in segment suggestions from multiple descriptors could indicate a region needing more detailed analysis.
- **Geometric Features and Energy Functions [26]:**
  - *Curvature:* High curvature regions often correspond to feature lines or detailed parts that are critical for accurate boundary definition. Mesh refinement can be triggered in such areas.[2]
  - *Homogeneity Energy (Ehom):* This measures how well a mesh element (or a small patch of elements) is contained within a single material type or a single segment. Adaptation aims to modify the mesh to lower this energy, ideally aligning element boundaries with segment boundaries. $Ehom=1-H$, where H is the ratio of the maximum area within an element belonging to a single material phase to the total area of the element.[26]
  - *Shape Energy (Eshape):* This quantifies the quality of mesh element shapes (e.g., deviation from ideal equilateral triangles or squares). Poorly shaped elements can negatively impact the accuracy of descriptor computations or numerical stability. Adapting the mesh to improve element shapes can thus indirectly benefit segmentation.[26]
  - *Combined Energy:* These energies can be combined, e.g., $E=\alpha Ehom+(1-\alpha)Eshape$, where α is a tunable parameter allowing prioritization between homogeneity and shape quality.[26]
- **User-Supplied Rules / Descriptor-Driven Criteria:** Adaptation can be guided by explicit rules or by the values of computed shape descriptors. For example:
  - Refine mesh regions where the gradient of a descriptor (like SDF or HKS) is high, indicating a potential boundary.
  - Increase mesh density near detected salient features.

It is important to recognize that true adaptive *segmentation* extends beyond mere adaptive *meshing*. While adaptive meshing techniques (like node movement or element splitting/merging) provide the operational tools, the crucial distinction lies in the criteria used to deploy these tools. For segmentation, these criteria must be specifically designed to improve segment quality—such as boundary accuracy or region homogeneity with respect to semantic labels—rather than focusing solely on generic geometric error or element shape quality, although the latter can be a prerequisite for reliable descriptor computation. For instance, a sharp gradient in HKS values near a current segment boundary might trigger not only mesh refinement in that local area but also a re-evaluation of the boundary position by the segmentation algorithm itself.

B. Methodologies in Adaptive Segmentation

Several methodologies can be employed to achieve adaptive segmentation:

- **Traditional Mesh Adaptation Techniques (applied to segmentation context [26]):** These techniques, often originating from finite element analysis, modify the mesh structure.
  - **Node Movement ("Movers"):** Vertex positions are adjusted to improve element quality or to align mesh edges with identified features or segment boundaries. Examples include:
    - *Anneal:* Randomly moves nodes, accepting moves that lower an effective energy (e.g., combined shape and homogeneity energy).
    - *Smooth:* Repositions each node to the average position of its neighbors.
    - *Snap Nodes/Snap Anneal:* Moves nodes to snap to nearby pixel boundaries (in image-based meshing) or feature lines.
    - *Relax:* Applies fictitious forces based on element energies to improve homogeneity and shape.
  - **Element Splitting/Merging ("Splitters" / Coarsening):**
    - *Refinement (Splitting):* Elements are subdivided in regions requiring more detail for accurate segmentation, where boundaries are uncertain, or where descriptor values change rapidly. Examples include bisecting edges or splitting elements at feature interfaces (Snap Refine).[26]
    - *Coarsening (Merging):* Elements are merged in smooth, homogeneous regions already confidently assigned to a segment, or where detail is unnecessary for the current segmentation task. This helps to reduce computational cost.[25]
- **Region-Based Refinement/Growth:**
  - The Reeb graph-based segmentation approach described in [2] employs an adaptive region-growing process. This process is guided by the topological structure of the Reeb graph but also considers mesh vertex connectivity to ensure contiguous segments, adapting the growth based on these dual constraints.
  - Another example is the adaptive mesh segmentation for dynamic delamination

modeling in composites.[27] Here, the mesh is adaptively segmented by adding new nodes "on-the-fly" to explicitly model discontinuities (cracks) in the displacement field. Cohesive segments are introduced between these new nodes, and their behavior is controlled by an energy criterion.

- **Advanced and Learning-Based Approaches:**
  - **Graph Transformers (e.g., MeshFormer [28]):** This represents a newer class of methods leveraging deep learning for adaptive segmentation, particularly for high-resolution meshes. MeshFormer addresses challenges such as large data and model sizes, and the difficulty of extracting sufficient semantic meaning from high-resolution data. Its key components include:
    1. A boundary-preserving simplification algorithm to reduce data size while protecting important segmentation boundaries.
    2. A Ricci flow-based clustering algorithm to construct hierarchical structures of the mesh. This allows for global receptive fields with fewer convolution layers, reducing model size.
    3. Graph attention convolutions applied within these hierarchical structures to effectively mine semantic meanings across different resolutions. This approach signifies a move towards learning-based adaptive strategies that inherently build and utilize hierarchical representations and adapt the model architecture or focus based on learned features.
  - **Multigrid-like Approaches:** While [29] primarily discusses solving PDEs using multiple predefined static refinement levels (a geometric multigrid method), the core concept of operating on different levels of detail and transferring information between them (restriction and prolongation) has parallels in hierarchical segmentation. A truly adaptive multigrid method for segmentation would dynamically adjust the refinement levels and locations based on segmentation-specific error estimators or confidence measures.

C. Benefits and Challenges in Practice

Adaptive segmentation offers significant potential but also comes with practical challenges.
Advantages [24]:

- **Computational Savings:** By concentrating computational effort and mesh density only in critical regions, adaptive methods can significantly reduce overall resource consumption compared to using uniformly fine meshes.
- **Storage Savings:** If coarsening is effectively employed in smooth or unimportant regions, the resulting meshes can be smaller, leading to reduced storage requirements.
- **Improved Accuracy:** Enhanced resolution in areas crucial for segmentation (e.g., complex boundaries, regions with fine details) can lead to more precise and geometrically accurate segmentations.
- **Handling Complex Geometries/Transient Events:** The ability to adapt dynamically allows these methods to better handle shapes with intricate details or scenarios where features relevant to segmentation evolve (e.g., in time-varying simulations or interactive

segmentation).

- **Less *a priori* Knowledge Required:** Compared to manually pre-tuning a static mesh for optimal segmentation, adaptive approaches can adjust based on the evolving data or intermediate segmentation results, requiring less detailed initial knowledge of where refinement is needed.

Challenges and Considerations [25]:

- **Initial Mesh Quality:** This is a critical factor. The adaptive process relies on the initial mesh being good enough to capture the essential characteristics of the geometry and flow field (in CFD contexts) or the fundamental shape features relevant for segmentation.[25] If the initial mesh is too coarse, especially near curved boundaries, sharp corners, or regions of high gradient, the adaptation process might fail to recover these features or could even degrade the mesh quality.[25]
- **Adaptation Criteria Design:** Defining robust and effective criteria that reliably trigger adaptation to improve *segmentation quality* (not just geometric accuracy or PDE solution error) is a non-trivial research problem. The criteria must be sensitive to features that matter for the specific segmentation task.
- **Convergence and Stability:** Ensuring that the adaptive process converges to a stable and improved segmentation, without oscillating or leading to runaway refinement, can be challenging. If adaptation is driven by an intermediate solution (e.g., a preliminary segmentation), that solution should ideally be well-converged before adaptation is applied.[25]
- **Parameter Sensitivity:** Adaptive methods often introduce their own set of parameters (e.g., thresholds for refinement/coarsening, weights for different energy terms like α in $E=\alpha Ehom+(1-\alpha)Eshape$ [26], error tolerances). Tuning these parameters for optimal performance on diverse datasets can be complex.
- **Implementation Complexity:** Managing dynamically changing mesh structures, connectivity information, and the logic for adaptation is significantly more complex than working with static meshes.
- **Cost of Adaptation:** The process of evaluating adaptation criteria, deciding where and how to adapt, and then performing the mesh modifications (e.g., splitting/merging elements, re-computing neighbors) incurs its own computational overhead. This cost must be balanced against the savings achieved.
- **High-Resolution Challenges [28]:** For very high-resolution meshes, issues like unstable mesh quality (e.g., non-manifold configurations, self-intersections) can hinder methods based on standard convolutions or geometric operations. Boundary-preserving simplification techniques become particularly important in such cases.
- **Disturbances from Adaptation [27]:** In methods that introduce new elements or significantly alter local topology during adaptation (e.g., adding cohesive segments in fracture modeling), care must be taken to minimize disturbances to surrounding fields or structures. Abrupt changes can cause numerical oscillations or inaccuracies if not handled with appropriate re-mapping or smoothing strategies.

A noteworthy consideration is the potential for a feedback loop: descriptors can guide adaptive meshing, and the refined mesh can, in turn, facilitate more accurate descriptor computation, ultimately leading to improved segmentation. For example, if initial descriptors computed on a coarse mesh suggest a potential boundary, the mesh can be refined in that area. Re-computing the descriptors on this locally refined mesh might provide clearer, more precise values, which then allow the segmentation algorithm to delineate the boundary with higher accuracy. However, managing such an iterative loop to ensure convergence, avoid excessive computational cost, and prevent instability (e.g., over-refinement in noisy areas) is a significant research challenge. Hierarchical approaches, like that of MeshFormer [28] which operates at different scales, represent one strategy to manage this complexity.

V. Synergies, Comparisons, and Evaluation Strategies

Understanding how different descriptor methods and adaptive techniques relate to each other, and how their performance can be objectively measured, is crucial for advancing the field of 3D mesh segmentation.

A. Pairwise Harmonics in the Landscape of Descriptors

The Pairwise Harmonics method [7] carves out a unique niche among shape descriptors due to its foundational pairwise analysis.

Comparative Discussion:

- **Global vs. Local Perspective:** Pairwise Harmonics inherently adopt a global perspective by considering the relationship and the connecting geometry between two points, even if these points are distant on the mesh.[7] This contrasts sharply with purely local descriptors like curvature [18], which only consider an infinitesimal neighborhood, and also differs from point-wise global descriptors like HKS/WKS [3], which, while influenced by global geometry through the Laplacian spectrum, are still defined at individual points.
- **Discriminative Power:** The information captured by the length (R) and average endpoint distance (D) of isocurves from pairwise harmonic fields [7] offers a distinct way to characterize shape. This can be compared to:
    - *Spectral methods:* Information from eigenvalues (frequencies) and eigenvectors (modes of variation).
    - *SDF:* Volumetric thickness measures.
    - *HKS/WKS:* Signatures of heat/wave diffusion patterns. The specific geometric nuances each type of descriptor is most sensitive to will vary, impacting their ability to distinguish different types of shape features relevant for segmentation. For example, pairwise harmonics might excel at identifying constrictions or symmetries between chosen point pairs.
- **Scale Handling:** Pairwise Harmonics claims an inherent multiscale analysis capability because the measurements are constrained along local geodesic paths, revealing structures at different scales for different point pairs without explicit scale parameters.[7] This can be contrasted with HKS and WKS, which use explicit time (t) or energy (E)

parameters to probe different scales [3], or with methods that require constructing explicit multi-resolution mesh hierarchies.

- **Robustness:** The reliance on harmonic functions, which are solutions to an elliptic PDE, endows Pairwise Harmonics with good robustness to mesh tessellation irregularities and surface noise.[7] This is comparable to the robustness achieved by the offset-SDF method [15] or the isometry invariance of HKS/WKS.
- **Suitability for Specific Segmentation Tasks:**
    - *Pairwise Harmonics:* Particularly strong for segmenting articulated structures (due to rigidity analysis in its segmentation pipeline) and symmetry-aware tasks.[7]
    - *Spectral Methods:* Well-suited for achieving globally consistent cuts and identifying intrinsic modes of the shape, often leading to "natural" partitions.[12]
    - *SDF:* Useful for segmentations based on thickness or volumetric properties, such as separating thin protrusions from thicker body parts.[14]
    - *HKS/WKS:* Effective for segmenting non-rigid or deformable shapes due to their isometry invariance, and for identifying multi-scale features.[3]
    - *Curvature/Saliency Methods:* Good for quickly identifying sharp features and boundaries but are primarily local in nature.[2]
    - *Reeb Graphs:* Excel at capturing the high-level topological structure of parts, such as limbs, handles, or tunnels.[21]

B. Integrating Descriptors with Adaptive Segmentation

A promising direction is the synergistic integration of sophisticated shape descriptors with adaptive segmentation methodologies. Descriptors can provide the semantic or geometric cues necessary to guide the adaptation process intelligently.
Guiding Adaptation:
Descriptor information can serve as the primary criteria for triggering and directing adaptive processes:

- **Boundary Refinement:** High gradients in descriptor values (e.g., a sharp change in SDF, HKS, or WKS values across a small region) near a currently estimated segment boundary can signal the need for local mesh refinement to capture the boundary more accurately.
- **Region of Interest Identification:** Saliency maps derived from geometric descriptors (e.g., curvature, HKS maxima) can identify regions of high geometric complexity or importance, which can then be prioritized for finer meshing or more detailed analysis by the segmentation algorithm.
- **Confidence-Based Adaptation:** If a segmentation algorithm produces low confidence scores for the labels of certain mesh elements (perhaps based on the ambiguity or similarity of their descriptor values to multiple segment profiles), these regions can be targeted for adaptive refinement to gather more information.

Hybrid Approaches:

- **Coarse-to-Fine Segmentation:** A global descriptor method (like Pairwise Harmonics, spectral methods, or Reeb graphs) could be used to obtain an initial, coarse segmentation of the shape or to identify major regions of interest. Subsequently, local descriptors (e.g., curvature, HKS/WKS values evaluated near the coarse boundaries) or error metrics could guide an adaptive meshing process to refine these boundaries with higher precision.

- **Improved Descriptor Computation:** There exists a potential feedback loop where adaptive meshing can improve the accuracy of the descriptors themselves. For instance, computing harmonic functions (for Pairwise Harmonics) or solving for Laplacian eigenfunctions (for HKS/WKS) on a mesh that has been adaptively refined to better capture the underlying geometry might lead to more accurate descriptor values. These improved descriptors could then lead to a better subsequent segmentation (related to the earlier discussion on the feedback loop in adaptive segmentation).

C. Overview of Evaluation Context for Mesh Segmentation

Objectively evaluating and comparing the performance of different 3D mesh segmentation algorithms necessitates the use of standardized benchmark datasets and quantitative evaluation metrics.[6]

Importance of Benchmarking:

Benchmarking provides a common ground for assessing the strengths and weaknesses of various methods, tracking progress in the field, and identifying areas for future improvement. Without standardized evaluation, comparisons become subjective and unreliable.

Common Benchmark Datasets 6:

Several datasets are commonly used in the literature, catering to different types of segmentation tasks:

- **Princeton Segmentation Benchmark (PSB):** A widely adopted dataset for evaluating part-based segmentation of single objects. It provides ground-truth segmentations for a variety of object categories.
- **SHREC (Shape Retrieval Contest):** This annual contest often includes tracks focused on 3D mesh segmentation (e.g., SHREC'12 [6]), providing new datasets and evaluation protocols.
- **ScanNet v2, Matterport3D, S3DIS (Stanford Large-scale 3D Indoor Spaces):** These are large-scale datasets of indoor scenes, typically captured with RGB-D sensors. They provide 3D mesh or point cloud data with dense semantic annotations (e.g., wall, floor, chair, table) and are extensively used for training and evaluating deep learning-based semantic segmentation methods.[30]
- **SUM (Semantic Urban Meshes), Hessigheim 3D (H3D):** These datasets provide urban-scale outdoor environments with semantic labels for elements like buildings, terrain, vegetation, and vehicles, often used for remote sensing and urban modeling applications.[30]

Standard Evaluation Metrics 6:

A variety of metrics are employed to quantify different aspects of segmentation quality:

- **Region Overlap/Similarity:**
  - **Intersection over Union (IoU) / Jaccard Index:** For each class or segment, IoU =

(Area of Overlap) / (Area of Union) between the predicted segment and the ground truth segment. The **mean IoU (mIoU)**, averaged over all classes, is a very common summary metric, especially in semantic segmentation.[30] Often, face areas are used in the computation for mesh-based IoU.[30]
  - **Dice Coefficient (F1-score):** Defined as 2×(Precision+Recall)(Precision×Recall) for segment overlap, it is equivalent to 2×(Predicted Size + Ground Truth Size)Overlap. It is also frequently reported.[30]
- **Boundary Accuracy:**
  - **Boundary Displacement Error (BDE):** Measures the average distance between the predicted segment boundaries and the ground truth boundaries.
  - **Boundary F1-score (BF1) / Boundary Precision (BP) / Boundary Recall (BR):** These metrics specifically evaluate how well the predicted boundaries align with the ground truth boundaries, considering a tolerance distance.[30]
- **Information Theory Metrics:**
  - **Variation of Information (VI):** Measures the amount of information lost or gained when transforming one clustering (segmentation) into another. Lower values indicate higher similarity.
  - **Rand Index (RI) and Adjusted Rand Index (ARI):** These metrics measure the similarity between two data clusterings by considering all pairs of elements and counting pairs that are assigned to the same or different clusters in both segmentations. ARI corrects for chance agreement.
- **Per-Class and Overall Accuracy:**
  - **Overall Accuracy (OA):** The percentage of correctly classified mesh elements (vertices or faces) over the total number of elements.[30]
  - **Mean per-class Accuracy (mAcc):** The average of accuracies computed for each individual class.[30]

The choice of specific metrics often depends on the nature of the segmentation task (e.g., instance segmentation vs. semantic segmentation vs. part-based decomposition) and what aspects of quality are most critical (e.g., precise boundary localization vs. correct overall part identification, penalties for over- vs. under-segmentation).

It is crucial to recognize that the choice of benchmark dataset and evaluation metric significantly influences the perceived performance and ranking of segmentation algorithms. A method that excels on a dataset of articulated objects evaluated with part-based metrics like the Rand Index (e.g., the PSB) might not perform as well on dense urban scene datasets evaluated with per-face semantic IoU (e.g., SUM or ScanNet). This context-dependency underscores the importance for researchers and practitioners to select benchmarks and metrics that closely align with their specific application goals and the types of shapes or scenes they intend to segment. "Good performance" is, therefore, a relative term, heavily conditioned by the evaluation framework.

VI. Strategic Recommendations and Future Outlook

Based on the comprehensive analysis of Pairwise Harmonics, other descriptor methods, and adaptive segmentation techniques, several strategic recommendations can be made for leveraging these approaches and for future research endeavors.

A. Leveraging User's "Pairwise Harmonics" Implementation

Given the user's existing implementation of "Pairwise Harmonics for Shape Analysis" [7], the following strategies are recommended:

- **Focus on Strengths:** The implementation should be primarily targeted at segmentation tasks where the method's inherent strengths are most advantageous. These include the segmentation of **articulated objects** (e.g., characters, animals), shapes exhibiting clear **symmetries**, and scenarios requiring high **robustness to noise and pose variations**.[7] The rigidity analysis component of its segmentation pipeline is particularly well-suited for these cases.
- **Systematic Parameter Exploration:** Achieving optimal results will necessitate a systematic exploration of the key parameters within the segmentation algorithm derived from.[7] This includes:
  - *Sampling density for point pairs:* Evaluate the trade-off between computational cost and the ability to capture all relevant structural details.
  - *Rigidity thresholds:* Fine-tune the threshold for classifying skeletal nodes as rigid or nonrigid, as this directly impacts initial part decomposition.
  - *Merging criteria:* Experiment with different parameters (e.g., volume thresholds, geometric similarity measures) for the segmentation refinement stage.
- **Addressing Limitations:** If the known limitation regarding components lacking a well-defined principal axis [7] proves problematic for target applications, consider hybrid approaches. For instance, another segmentation method could be used to pre-segment the object, with Pairwise Harmonics then applied to refine segments or to analyze specific sub-parts. Alternatively, regions where Pairwise Harmonics struggle could be post-processed using a different technique.

B. Selecting or Combining Other Descriptor Methods

The choice of alternative or complementary descriptor methods should be guided by the specific requirements of the application:

- **Based on Application Needs:**
  - For **non-rigid, deformable shapes** where isometry invariance is paramount, **HKS or WKS** are strong candidates due to their spectral nature and proven robustness to such deformations.[3]
  - If the objects have a clear **topological structure** (e.g., distinct limbs, handles, holes), **Reeb graphs** can provide an excellent high-level decomposition that captures this skeletal abstraction.[21]
  - When dealing with noisy or incomplete data and requiring extreme robustness, the **offset-SDF method** offers a reliable way to capture volumetric properties.[14]
  - For achieving **globally consistent cuts** based on intrinsic shape modes, traditional **spectral methods** based on the Laplacian eigenstructure remain

powerful.[12]

- **Hybrid Strategies:**
  - **Coarse-to-Fine Segmentation:** Employ a global method like Pairwise Harmonics, spectral analysis, or Reeb graphs to obtain an initial, coarse segmentation. Then, use local methods (e.g., curvature analysis, or HKS/WKS values evaluated near the coarse boundaries) or an adaptive meshing strategy to refine these boundaries with greater precision.
  - **Descriptor Fusion:** For tasks requiring very rich feature representations, consider combining features from multiple descriptor types (e.g., concatenating feature vectors from Pairwise Harmonics and HKS). This approach, noted in [8] as being widely adopted, can potentially capture a more comprehensive set of shape characteristics, though it increases computational complexity and requires careful feature selection or weighting.

C. Incorporating Adaptive Techniques

Adaptive segmentation techniques can offer significant benefits, particularly under certain conditions:

- **When to Consider Adaptation:** Adaptive methods are most valuable for segmenting **complex shapes**, meshes exhibiting **highly variable detail density** across their surface, or when **high precision is required in specific critical areas** (e.g., intricate boundaries) while needing to manage overall computational cost.
- **Guiding Adaptation with Descriptors:** The output of the chosen shape descriptor(s) can serve as the primary input for guiding the adaptive process. For example:
  - Mesh resolution can be increased in regions where descriptor values change rapidly (e.g., high gradient of SDF or HKS), indicating a potential feature or boundary.
  - If the segmentation algorithm exhibits low confidence in its labeling of certain mesh elements (perhaps due to ambiguous descriptor values), these areas can be targeted for adaptive refinement to gather more distinguishing information.
- **Starting Simple:** For users new to adaptive segmentation, it is advisable to begin with conceptually simpler techniques. This might involve refining the mesh based on local curvature values or proximity to estimated boundaries before attempting more complex, learning-based adaptive models like MeshFormer [28], which require significant data and expertise in deep learning.

D. Potential Avenues for Further Research and Advanced Development

The field of 3D mesh segmentation continues to evolve, with several promising avenues for future research:

- **Learned Pairwise Descriptors:** Extending the philosophy of Pairwise Harmonics or Pairwise 3D Shape Context, deep learning approaches could be explored to learn optimal pairwise features directly from data. This could lead to descriptors that are more discriminative and better tailored to specific types of shapes or segmentation tasks.

- **Automatic Parameter Tuning:** A persistent challenge in many descriptor-based and adaptive methods is the need for manual parameter tuning (e.g., k for spectral clustering, t for HKS, α for energy-based adaptation). Research into methods that can automatically determine optimal parameters based on the characteristics of the input mesh or the desired segmentation outcome would significantly enhance usability and robustness.
- **Tighter Integration of Descriptors and Adaptive Solvers:** Developing frameworks where descriptor computation and adaptive mesh/segmentation refinement are more deeply intertwined within an iterative optimization loop could lead to more powerful and efficient systems. This involves creating a feedback mechanism where improved meshes lead to better descriptors, which in turn guide further adaptation and segmentation.
- **Explainable AI (XAI) for Segmentation:** As learning-based methods, particularly deep learning, become more prevalent in mesh segmentation, there is a growing need for techniques that can provide insights into *why* these models produce certain segmentations. XAI methods could increase the trustworthiness and interpretability of these complex algorithms.

The overarching trend in mesh segmentation appears to be moving towards hybrid approaches. No single descriptor or method is universally optimal for all types of shapes, noise levels, or segmentation goals. Consequently, the future likely lies in developing flexible frameworks that can intelligently combine the strengths of different descriptor families (e.g., global understanding from spectral or harmonic methods with local detail from curvature or diffusion signatures) and leverage adaptive techniques to focus computational effort where it is most effective. The primary challenge will be in creating principled ways to fuse diverse sources of information and manage the complex interactions between different algorithmic components.

VII. Conclusion

A. Recapitulation of Key Findings

This report has undertaken a detailed examination of descriptor-based and adaptive techniques for 3D mesh segmentation, with an initial focus on the "Pairwise Harmonics for Shape Analysis" method. The analysis reveals that Pairwise Harmonics offers a unique global pairwise perspective, demonstrating robustness to noise and pose variations, and proving particularly effective for segmenting articulated structures and shapes with symmetries.[7] However, its efficacy can be dependent on components having well-defined principal axes and careful sparse sampling of point pairs.

Beyond this specific method, a diverse landscape of other prominent descriptor families has been surveyed. Spectral descriptors, derived from the eigenstructures of mesh Laplacians, provide powerful tools for capturing global intrinsic shape modes and enabling consistent cuts.[12] Volumetric descriptors like the Shape Diameter Function (SDF), especially its robust offset-based variant, are valuable for segmentations guided by local thickness.[14] Point-based diffusion signatures such as the Heat Kernel Signature (HKS) and Wave Kernel Signature (WKS) offer multi-scale, isometry-invariant features crucial for analyzing non-rigid shapes.[3] Geometric saliency and curvature-driven methods provide intuitive and often fast ways to detect sharp features, albeit with a local focus.[18] Topological descriptors, exemplified by Reeb graphs, excel at abstracting high-level part structures.[21]

Furthermore, the principles and methodologies of adaptive mesh segmentation have been explored. These techniques, by dynamically adjusting mesh resolution or the segmentation process itself based on evolving features or confidence metrics, hold significant potential for enhancing segmentation quality and computational efficiency.[24] However, their practical implementation involves challenges related to initial mesh quality, the design of effective adaptation criteria, and overall algorithmic complexity.[25]

B. Synthesis and Final Thoughts

The selection of an appropriate 3D mesh segmentation strategy is unequivocally contingent upon the specific application context, the inherent characteristics of the input meshes (e.g., complexity, noise levels, presence of articulation), the desired granularity and semantic meaning of the segments, and the available computational resources. There is no universally superior method; rather, each approach presents a unique set of trade-offs.

A profound understanding of the underlying principles, strengths, and limitations of different descriptor families and adaptive techniques is paramount for making informed decisions in research and development. The "Pairwise Harmonics" approach provides a valuable tool, particularly for its targeted strengths, but its utility can be further amplified when considered within the broader ecosystem of available methods. Hybrid strategies, combining global understanding from one class of descriptors with local refinement from another, or integrating descriptor-driven cues into adaptive frameworks, appear to be a particularly promising direction.

The field of 3D mesh segmentation is continuously advancing, driven by the increasing complexity and diversity of 3D data and the escalating demands of downstream applications. The trend points towards more robust, flexible, and increasingly data-driven (learning-based) approaches that can intelligently navigate these challenges. It is hoped that the comprehensive analysis and strategic recommendations provided in this report will equip the user with the necessary insights to effectively advance their specific mesh segmentation objectives.

Works cited

1. Deep Learning Based 3D Segmentation: A Survey - arXiv, accessed on May 7, 2025, https://arxiv.org/html/2103.05423v5
2. Flexible Mesh Segmentation via Reeb Graph Representation of Geometrical and Topological Features - arXiv, accessed on May 7, 2025, https://arxiv.org/html/2412.05335v2
3. geometry.stanford.edu, accessed on May 7, 2025,

https://geometry.stanford.edu/lgl_2024/papers/socg-pbsds-10/socg-pbsds-10.pdf
4. Discrete Dierential-Geome try Operators for Triangulated 2 ..., accessed on May 7, 2025, https://www.researchgate.net/publication/242361656_Discrete_Dierential-Geome_try_Operators_for_Triangulated_2-Manifolds
5. Part-Based Mesh Segmentation: A Survey: Mesh Segmentation ..., accessed on May 7, 2025, https://www.researchgate.net/publication/322868499_Part-Based_Mesh_Segmentation_A_Survey_Mesh_Segmentation
6. A comprehensive overview of methodologies and performance ..., accessed on May 7, 2025, https://www.researchgate.net/publication/271141076_A_comprehensive_overview_of_methodologies_and_performance_evaluation_frameworks_in_3D_mesh_segmentation
7. web.engr.oregonstate.edu, accessed on May 7, 2025, https://web.engr.oregonstate.edu/~zhange/images/pairwise_harmonics.pdf
8. 3D Shape Segmentation with Potential Consistency Mining and Enhancement, accessed on May 7, 2025, https://nia.nbt.edu.cn/PDF/45_3D_Shape_Segmentation_with_Potential_Consistency_Mining_and_Enhancement.pdf
9. 3D Representation Methods: A Survey - arXiv, accessed on May 7, 2025, https://arxiv.org/html/2410.06475v1
10. www.cse.iitb.ac.in, accessed on May 7, 2025, https://www.cse.iitb.ac.in/~cs749/spr2017/handouts/levy_spectral.pdf
11. Spectral mesh processing | Request PDF - ResearchGate, accessed on May 7, 2025, https://www.researchgate.net/publication/220720161_Spectral_mesh_processing
12. (PDF) Spectral Mesh Processing. - ResearchGate, accessed on May 7, 2025, https://www.researchgate.net/publication/220507645_Spectral_Mesh_Processing
13. Spectral 3D Mesh Segmentation with a Novel Single Segmentation ..., accessed on May 7, 2025, https://www.researchgate.net/publication/261715250_Spectral_3D_Mesh_Segmentation_with_a_Novel_Single_Segmentation_Field
14. Fast and robust shape diameter function - PMC, accessed on May 7, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC5786294/
15. Fast and robust shape diameter function - ResearchGate, accessed on May 7, 2025, https://www.researchgate.net/publication/322727563_Fast_and_robust_shape_diameter_function
16. openaccess.thecvf.com, accessed on May 7, 2025, https://openaccess.thecvf.com/content_cvpr_2015/papers/Fang_3D_Deep_Shape_2015_CVPR_paper.pdf
17. etheses.whiterose.ac.uk, accessed on May 7, 2025, https://etheses.whiterose.ac.uk/id/eprint/18036/8/amendments_limberger_fa_Spe

ctral%20Signatures%20for%20Non-rigid%203D%20Shape%20Retrieval.pdf
18. arxiv.org, accessed on May 7, 2025, https://arxiv.org/html/2412.05335v1#:~:text=Surface%2Dbased%20mesh%20segmentation%20algorithms,growing%2C%20clustering%20and%20boundary%20segmentation.
19. egusphere.copernicus.org, accessed on May 7, 2025, https://egusphere.copernicus.org/preprints/2025/egusphere-2025-1110/egusphere-2025-1110.pdf
20. (PDF) Curvature-based pebble segmentation for reconstructed ..., accessed on May 7, 2025, https://www.researchgate.net/publication/390340533_Curvature-based_pebble_segmentation_for_reconstructed_surface_meshes
21. Flexible Mesh Segmentation through Integration of Geometric and Topological Features of Reeb Graphs - arXiv, accessed on May 7, 2025, https://arxiv.org/html/2412.05335v1
22. arxiv.org, accessed on May 7, 2025, https://arxiv.org/html/2412.05335v2#:~:text=The%20algorithm%20consists%20of%20three,region%2Dgrowth%20process%20that%20takes
23. Pairwise Three-Dimensional Shape Context for Partial Object Matching and Retrieval on Mobile Laser Scanning Data - ResearchGate, accessed on May 7, 2025, https://www.researchgate.net/publication/260712540_Pairwise_Three-Dimensional_Shape_Context_for_Partial_Object_Matching_and_Retrieval_on_Mobile_Laser_Scanning_Data
24. Adaptive mesh refinement - Wikipedia, accessed on May 7, 2025, https://en.wikipedia.org/wiki/Adaptive_mesh_refinement
25. Accelerate CFD Simulation with Adaptive Mesh Refinement: Best ..., accessed on May 7, 2025, https://blog.ozeninc.com/resources/best-practice-for-adaptive-mesh-refinement-in-ansys-fluent
26. www.ctcms.nist.gov, accessed on May 7, 2025, https://www.ctcms.nist.gov/~rlua/FE/meshing.pdf
27. Adaptive Mesh Segmentation for Modelling Dynamic Delamination Initiation and Propagation in Thick Composite Laminates - Dynalook, accessed on May 7, 2025, https://www.dynalook.com/conferences/12th-european-ls-dyna-conference-2019/fiber-reinforced-polymers/selvaraj_university-of-bristol.pdf
28. MeshFormer: High-resolution Mesh Segmentation with Graph Transformer - Eurographics Association, accessed on May 7, 2025, https://diglib.eg.org/bitstream/handle/10.1111/cgf14655/v41i7pp037-049.pdf
29. Algorithm for Adaptive Mesh Refinement - Computational Science Stack Exchange, accessed on May 7, 2025, https://scicomp.stackexchange.com/questions/24988/algorithm-for-adaptive-mesh-refinement
30. uwaterloo.ca, accessed on May 7, 2025, https://uwaterloo.ca/geospatial-intelligence/sites/default/files/uploads/documents

/deep-learning-based-semantic-segmentation-of-urban-scale-3d-meshes-in.pd
f